

```

# train_model_full_set.py
# -----
# Trains a ResNet50 image classification model on refund data.
# Logs training and validation metrics using MLflow for tracking.
# Saves the final model as refund_classifier_final.pt for the API.
# Designed for reproducible and scalable model development.
# -----

import os
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, models, transforms
from torchvision.models import resnet50, ResNet50_Weights
from torch.utils.data import DataLoader
from tqdm import tqdm
import mlflow
import mlflow.pytorch
from collections import Counter

# --- CONFIGURATION ---
EPOCHS = 30
LR = 0.001
BATCH_SIZE = 32
MODEL_PATH = "refund_classifier_final.pt"
DATA_DIR = "/content/data/full"

# --- TRANSFORMS ---
train_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(15),
    transforms.ColorJitter(brightness=0.2, contrast=0.2),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
                          [0.229, 0.224, 0.225])
])
val_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
                          [0.229, 0.224, 0.225])
])

# --- DATA LOADERS ---
train_data = datasets.ImageFolder(os.path.join(DATA_DIR, 'train'), transform=train_transform)
val_data = datasets.ImageFolder(os.path.join(DATA_DIR, 'val'), transform=val_transform)
train_loader = DataLoader(train_data, batch_size=BATCH_SIZE, shuffle=True)
val_loader = DataLoader(val_data, batch_size=BATCH_SIZE)

print("Train class distribution:")
train_labels = [label for _, label in train_data.imgs]
print(Counter(train_labels))

print("Val class distribution:")

```

```

val_labels = [label for _, label in val_data.imgs]
print(Counter(val_labels))

# --- MODEL SETUP ---
model = resnet50(weights=ResNet50_Weights.DEFAULT)
model.fc = nn.Sequential(
    nn.Dropout(0.3),
    nn.Linear(model.fc.in_features, len(train_data.classes))
)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=LR, momentum=0.9)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', patience=3, factor=0.5, verbose=True)

# --- MLflow Tracking ---
mlflow.set_experiment("Refund_Classifier")
with mlflow.start_run():
    mlflow.log_params({
        "epochs": EPOCHS,
        "learning_rate": LR,
        "batch_size": BATCH_SIZE,
        "optimizer": "SGD",
        "model": "resnet50"
    })

    best_val_acc = 0
    patience = 5
    trigger = 0

    for epoch in range(EPOCHS):
        model.train()
        running_loss = 0.0
        correct = 0
        total = 0
        for inputs, labels in tqdm(train_loader, desc=f"Epoch {epoch+1}/{EPOCHS}"):
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()
            _, predicted = outputs.max(1)
            total += labels.size(0)
            correct += predicted.eq(labels).sum().item()

        train_loss = running_loss / len(train_loader)
        train_acc = correct / total

    # Evaluate
    model.eval()
    val_loss = 0.0
    val_correct = 0
    val_total = 0
    with torch.no_grad():
        for inputs, labels in val_loader:
            outputs = model(inputs)

```

```

        loss = criterion(outputs, labels)
        val_loss += loss.item()
        _, predicted = outputs.max(1)
        val_total += labels.size(0)
        val_correct += predicted.eq(labels).sum().item()

val_loss /= len(val_loader)
val_acc = val_correct / val_total

print(f"Epoch {epoch+1}/{EPOCHS} - Train Acc: {train_acc:.4f}, Val Acc: {val_acc:.4f}")

# Log to MLflow
mlflow.log_metric("train_loss", train_loss, step=epoch)
mlflow.log_metric("train_acc", train_acc, step=epoch)
mlflow.log_metric("val_loss", val_loss, step=epoch)
mlflow.log_metric("val_acc", val_acc, step=epoch)

# Learning rate adjustment
scheduler.step(val_loss)

# Early stopping
if val_acc > best_val_acc:
    best_val_acc = val_acc
    trigger = 0
else:
    trigger += 1

if trigger >= patience:
    print("Early stopping triggered.")
    break

# Save model + log artifact
torch.save(model, MODEL_PATH)
mlflow.pytorch.log_model(model, "model")

# Save to Google Drive
from google.colab import drive
drive.mount('/content/drive')
torch.save(model, "/content/drive/MyDrive/refund_classifier_final.pt")

# Download model directly to local PC
from google.colab import files
files.download(MODEL_PATH)

```

```

🔗 Train class distribution:
Counter({28: 340, 17: 325, 13: 147, 29: 140, 3: 105, 25: 91, 8: 68, 12: 62, 9: 52, 7: 46, 16: 42, 21: 34, 27: 31, 14: 23, 10: 22, 22: 20, 20: 18, 23: 16, 6: 14, 24: 11, 26: 11, 19: 10, 2: 9, 15: 9})
Val class distribution:
Counter({28: 89, 17: 75, 13: 39, 29: 32, 3: 24, 12: 20, 25: 20, 8: 16, 16: 14, 9: 13, 7: 10, 14: 10, 10: 8, 21: 8, 27: 8, 20: 6, 22: 5, 24: 5, 4: 4, 0: 3, 6: 3, 15: 3, 26: 3, 2: 2, 23: 2, 1: 1, 5: 1})
Downloading: "https://download.pytorch.org/models/resnet50-11ad3fa6.pth" to /root/.cache/torch/hub/checkpoints/resnet50-11ad3fa6.pth
100%|██████████| 97.8M/97.8M [00:00<00:00, 261MB/s]
/usr/local/lib/python3.11/dist-packages/torch/optim/lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.
  warnings.warn(
2025/07/15 00:15:44 INFO mlflow.tracking.fluent: Experiment with name 'Refund_Classifier' does not exist. Creating a new experiment.
Epoch 1/30: 100%|██████████| 53/53 [04:27<00:00, 5.04s/it]
Epoch 1/30 - Train Acc: 0.2202, Val Acc: 0.3607
Epoch 2/30: 100%|██████████| 53/53 [04:22<00:00, 4.95s/it]
Epoch 2/30 - Train Acc: 0.4458, Val Acc: 0.5808
Epoch 3/30: 100%|██████████| 53/53 [04:01<00:00, 4.56s/it]
Epoch 3/30 - Train Acc: 0.6170, Val Acc: 0.6417
Epoch 4/30: 100%|██████████| 53/53 [04:01<00:00, 4.57s/it]
Epoch 4/30 - Train Acc: 0.6942, Val Acc: 0.7002
Epoch 5/30: 100%|██████████| 53/53 [03:58<00:00, 4.50s/it]
Epoch 5/30 - Train Acc: 0.7439, Val Acc: 0.7447
Epoch 6/30: 100%|██████████| 53/53 [04:01<00:00, 4.55s/it]
Epoch 6/30 - Train Acc: 0.7666, Val Acc: 0.7822
Epoch 7/30: 100%|██████████| 53/53 [03:58<00:00, 4.49s/it]
Epoch 7/30 - Train Acc: 0.8061, Val Acc: 0.7963
Epoch 8/30: 100%|██████████| 53/53 [04:02<00:00, 4.58s/it]
Epoch 8/30 - Train Acc: 0.8193, Val Acc: 0.7963
Epoch 9/30: 100%|██████████| 53/53 [04:01<00:00, 4.55s/it]
Epoch 9/30 - Train Acc: 0.8384, Val Acc: 0.8220
Epoch 10/30: 100%|██████████| 53/53 [03:59<00:00, 4.52s/it]
Epoch 10/30 - Train Acc: 0.8540, Val Acc: 0.8197
Epoch 11/30: 100%|██████████| 53/53 [04:00<00:00, 4.54s/it]
Epoch 11/30 - Train Acc: 0.8642, Val Acc: 0.8150
Epoch 12/30: 100%|██████████| 53/53 [04:00<00:00, 4.55s/it]
Epoch 12/30 - Train Acc: 0.8827, Val Acc: 0.8407
Epoch 13/30: 100%|██████████| 53/53 [04:00<00:00, 4.54s/it]
Epoch 13/30 - Train Acc: 0.8995, Val Acc: 0.8501
Epoch 14/30: 100%|██████████| 53/53 [04:01<00:00, 4.56s/it]
Epoch 14/30 - Train Acc: 0.9096, Val Acc: 0.8431
Epoch 15/30: 100%|██████████| 53/53 [04:03<00:00, 4.60s/it]
Epoch 15/30 - Train Acc: 0.9204, Val Acc: 0.8431
Epoch 16/30: 100%|██████████| 53/53 [03:59<00:00, 4.53s/it]
Epoch 16/30 - Train Acc: 0.9276, Val Acc: 0.8454
Epoch 17/30: 100%|██████████| 53/53 [03:57<00:00, 4.48s/it]
Epoch 17/30 - Train Acc: 0.9234, Val Acc: 0.8384
Epoch 18/30: 100%|██████████| 53/53 [04:02<00:00, 4.58s/it]
2025/07/15 01:33:48 WARNING mlflow.models.model: `artifact_path` is deprecated. Please use `name` instead.
Epoch 18/30 - Train Acc: 0.9437, Val Acc: 0.8501
Early stopping triggered.
2025/07/15 01:33:48 WARNING mlflow.utils.requirements_utils: Found torch version (2.6.0+cu124) contains a local version label (+cu124). MLflow logged a pip requirement for this package as 'torch==2.6.0+cu124'
2025/07/15 01:33:56 WARNING mlflow.utils.requirements_utils: Found torchvision version (0.21.0+cu124) contains a local version label (+cu124). MLflow logged a pip requirement for this package as 'torchvision==0.21.0+cu124'
2025/07/15 01:33:56 WARNING mlflow.models.model: Model logged without a signature and input example. Please set `input_example` parameter when logging the model to auto infer the model signature.
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```

```

# visualize_training_results.py
# -----
# Script to extract and plot accuracy and loss

```

```

# from MLflow logs after training is complete.
# -----

import mlflow
import pandas as pd
import matplotlib.pyplot as plt

# Load latest MLflow run (now that mlruns is available)
client = mlflow.tracking.MlflowClient()
experiment = client.get_experiment_by_name("Refund_Classifier")
runs = client.search_runs(experiment.experiment_id)

# Just use the first available run
run = runs[0]
run_id = run.info.run_id

# Retrieve metrics
metrics = client.get_metric_history(run_id, "train_acc")
train_acc = [m.value for m in metrics]
train_steps = [m.step for m in metrics]

metrics = client.get_metric_history(run_id, "val_acc")
val_acc = [m.value for m in metrics]

metrics = client.get_metric_history(run_id, "train_loss")
train_loss = [m.value for m in metrics]

metrics = client.get_metric_history(run_id, "val_loss")
val_loss = [m.value for m in metrics]


# Plot
plt.figure(figsize=(10, 4))

plt.subplot(1, 2, 1)
plt.plot(train_steps, train_acc, label="Train Accuracy")
plt.plot(train_steps, val_acc, label="Val Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Accuracy Over Epochs")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(train_steps, train_loss, label="Train Loss")
plt.plot(train_steps, val_loss, label="Val Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Loss Over Epochs")
plt.legend()

plt.tight_layout()
plt.show()

```



```
print("Train acc:", train_acc)
print("Val acc:", val_acc)
```

```
Train acc: [0.22022740873728305, 0.44584081388390184, 0.6169958108916817, 0.6941950927588271, 0.743865948533812, 0.7666068222621185, 0.8061041292639138, 0.8192698982645122, 0.8384201077199281, 0.8575201077199281, 0.8765701077199281, 0.8956201077199281, 0.9146701077199281, 0.9337201077199281, 0.9527701077199281, 0.9718201077199281, 0.9908701077199281]
Val acc: [0.36065573770491804, 0.5807962529274004, 0.6416861826697893, 0.7002341920374707, 0.7447306791569087, 0.7822014051522248, 0.7962529274004684, 0.7962529274004684, 0.8220140515222483, 0.8192698982645122, 0.8192698982645122, 0.8192698982645122, 0.8192698982645122, 0.8192698982645122, 0.8192698982645122, 0.8192698982645122, 0.8192698982645122]
```

```
# confusion_matrix_heatmap.py
# -----
# Script to generate a normalized confusion matrix
# for the trained refund classification model using
# the validation set. Highlights true vs. predicted
# class performance with a heatmap.
# -----

import torch
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
import numpy as np
from torchvision.datasets import ImageFolder
from torch.utils.data import DataLoader
from torchvision import transforms

# --- Load final model ---
model = torch.load("refund_classifier_final.pt")
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
model.eval()

# --- Load validation data ---
val_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
```

```

        transforms.Normalize([0.485, 0.456, 0.406],
                              [0.229, 0.224, 0.225])
    ])

val_data = ImageFolder("data/full/val", transform=val_transform)
val_loader = DataLoader(val_data, batch_size=32, shuffle=False)
class_names = val_data.classes

# --- Predict on validation set ---
all_preds = []
all_labels = []

with torch.no_grad():
    for inputs, labels in val_loader:
        inputs = inputs.to(device)
        outputs = model(inputs)
        _, preds = torch.max(outputs, 1)
        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

# --- Compute normalized confusion matrix ---
cm = confusion_matrix(all_labels, all_preds, normalize="true")

# --- Plot confusion matrix as heatmap ---
plt.figure(figsize=(12, 10))
sns.heatmap(cm, annot=True, fmt=".2f", cmap="YlGnBu",
            xticklabels=class_names,
            yticklabels=class_names,
            cbar_kws={"label": "Proportion"})
plt.title("Normalized Confusion Matrix (Validation Set)")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.xticks(rotation=45, ha="right")
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()

```



True Label

Normalized Confusion Matrix (Validation Set)

